Where the Human Comes In

A structural human gate for high-stakes AI decisions

Project	Witness - a human-in-the-loop decision desk
Event	Agentika 2.0 (NaSCon'26, FAST NUCES Islamabad)
Theme	Human + AI: Better Together
Live service	https://alternate-echo-whale-5578-github-app.alternatefutures.ai
Version / build	0.4.0 / image 8c0478e
Date	17 October 2026

Witness screens applicants for a simplified, synthetic income-support programme. A language model extracts evidence; ordinary code decides the rules; an independent check audits the result; and a human officer makes the decision. The question this document answers is narrow and deliberate: **where, exactly, does the human come in, and what stops the machine from deciding without them?**

> *The short answer: the human is not a reviewer bolted on after the fact. The automated pipeline is structurally incapable of producing a denial. Its strongest possible output is **NEEDS_REVIEW**. **DENIED** exists only as the result of an explicit human action, recorded with the officer's name and a written reason in a tamper-evident ledger.*

1. The failure we design against

In benefits screening the costly errors are asymmetric. A wrongly denied applicant may lose food support; a wrongly approved one costs money. Worse, an automated denial is hard to contest because nobody is accountable for it. Two common designs fail here:

• **The autonomous scorer.** A model outputs approve/deny. It is fast, but a denial has no owner and no reviewable basis.

• **The cosmetic reviewer.** A person clicks 'confirm' on whatever the model produced. The human step exists on the diagram but has no authority and no information advantage.

Witness rejects both. The machine is allowed to be useful without being allowed to be final, and the human is given both the authority and the evidence needed to exercise it.

2. The rule: the human is a stage in the pipeline

Case state is a closed set. The automated pipeline can only move a case into VERIFIED or NEEDS_REVIEW. Both are pre-decision states. A final state requires a human decision.

State	Reached by	Final?
INTAKE	Case received	No
VERIFIED	Every rule PASS, every citation grounded, verifier OK	No - awaiting officer
NEEDS_REVIEW	A rule FAILed, a rule is UNKNOWN, or the verifier abstained	No - awaiting officer
APPROVED	Human officer approves, with a written reason	Yes
DENIED	Human officer denies, with a written reason	Yes
INFO_REQUESTED	Human officer asks for more evidence	No - case reopens

There is no code path in which the pipeline writes APPROVED or DENIED. This is enforced by the store's decision gate, not by a prompt instruction: only the human decision endpoint can set a final state, and it refuses to act on a case that is not in a decidable state.

3. The pipeline, stage by stage

Stage	What it does	What it may not do
Reader	Turns a text or PDF case packet into documents (PDF pages become citable units).	Invent evidence.
Extractor	Proposes a value plus the exact quote it came from; the quote is re-grounded verbatim.	Return a fact without a verbatim citation.
Rules engine	Ordinary Python compares grounded facts to published thresholds; each rule cites a clause.	Let the model decide a rule.
Verifier	Re-audits grounding, contradictions and undecided rules; may only ADD an abstention.	Clear a deterministic failure or approve a case.
Human officer	Approves, denies, or requests information - with a mandatory written reason.	Decide without an identity or a reason.
Audit ledger	Appends a SHA-256 hash-chained entry for every state change.	Rewrite history undetected.

Two properties matter more than the individual stages. First, the model's opinion can only make the outcome *more cautious*: the LLM verifier may add an abstention but never removes a rule failure. Second, every claimed fact must survive a verbatim check against its source document, so a hallucinated quote produces no citation and the verifier abstains.

4. Where the human acts

An officer sees the evidence - each fact next to the exact quoted span - the rule results, the verifier's checks, and any conflicting values side by side. They then choose one of three actions:

• **Approve** or **deny**: a final decision. A denial is possible *only* here.

• **Request more information**: the case becomes INFO_REQUESTED. New evidence is appended and the pipeline re-runs from the top; the case reopens with no decision attached.

The backend enforces the terms of the act, not just the UI: the officer name must be non-empty, the reason must contain real text, the action must be one of the three, and the case must be in VERIFIED or NEEDS_REVIEW. Duplicate decisions and invalid transitions are rejected. When an officer decides a case the verifier had abstained on, the ledger records **overrode_abstention: true**, so the override is visible rather than silent.

5. Abstention instead of guessing

The verifier is designed to refuse. It abstains, and forces NEEDS_REVIEW, when any of the following is true:

• a fact's citation quote does not appear **verbatim** in its source document;

• the same field carries **conflicting** values across documents;

• a rule cannot be decided because its evidence is missing, unreadable, or contradictory.

This is the opposite of a confident chatbot. An abstention is a successful outcome: it routes an uncertain case to the person who is allowed to resolve it, instead of manufacturing a plausible answer. In the live evaluation set, 3 of 5 seeded cases abstain for exactly these reasons.

6. Rules are code, not vibes

Seven clauses are evaluated deterministically. The model never decides a rule, and each result names the clause it implements, so a decision traces back to written policy.

Clause	Rule id	Plain rule	Threshold
1	AGE_ADULT	Applicant is 18 or older	age >= 18
2	RESIDENCY	Resident in a covered region	province in covered list
3	HOUSEHOLD_SIZE	Household size recorded	household_size >= 1
4	INCOME_LIMIT	Monthly income within limit	monthly_income_pkr <= 60,000
5	ASSET_CAR	No car declared	owns_car == false
6	ASSET_LAND	Land holding within limit	land_acres <= 2
7	BANK_BALANCE	Liquid savings within limit	bank_balance_pkr <= 50,000

Disclaimer. The rule set is simplified, synthetic and BISP-inspired. It is a demonstration and does not represent the official policy of BISP or of any government or agency. Thresholds are illustrative.

7. Accountability: a hash-chained ledger

Every state change is appended to an append-only ledger. Each entry stores the hash of the previous entry, so altering or removing a single field breaks verification at that point. **GET /api/audit/verify** recomputes the whole chain and reports the first break.

The ledger is also written **before** in-memory state changes. If the ledger cannot be written, the API returns 503 and the case is left unchanged. A decision can therefore never appear to succeed without a record. This matters because a verifying chain proves *integrity*, not *durability*; Witness is explicit about the difference.

A live demonstration endpoint, **POST /api/audit/tamper-demo**, tampers with an in-memory copy of the ledger, shows the chain breaking exactly at the forged entry, and then re-verifies the untouched ledger to show the real record is intact.

8. What the AI is not allowed to do

Capability	Automated pipeline	Human officer
Extract facts with citations	Yes	Reviews them
Decide a policy rule	No (code does)	May override with a reason
Produce ELIGIBLE recommendation	Yes	May accept
Produce a final denial	**Never**	Yes, with a written reason
Clear a verifier abstention	No	Yes, recorded as overrode_abstention
Write the audit entry	Yes	The entry names the officer

9. Evidence from the live service

Check	Live result (17 Oct 2026)
Startup evaluation gate	55 / 55 expectations met; all_passed = true
Grounding rate	0.971 (facts carrying a verbatim citation)
Abstain rate	0.6 (cases escalated instead of guessed)
No-auto-deny invariant	true (pre-human states: INTAKE, VERIFIED, NEEDS_REVIEW)
Audit chain	verifies on the live ledger
Deployed build	v0.4.0, image 8c0478e

10. Honest limits

• **The durable database is not live.** A managed Postgres service exists and the adapter is implemented, but the two services were placed on different decentralized providers, so the cluster-local name does not resolve. Witness degrades to the embedded SQLite ledger and reports the degradation explicitly at /api/adapters rather than pretending.

• **Confidential compute is not running.** Two confidential redeploys of a duplicate service failed with 'CVM failed' (no viable TEE offer at the time). No attestation claim is made.

• **SQLite survives a process restart, not a redeploy.** This is stated plainly because the ledger proves integrity, not persistence.

• **The policy is synthetic.** The rules are a demonstration of structure, not real eligibility law.

11. Verify it in sixty seconds

Open https://alternate-echo-whale-5578-github-app.alternatefutures.ai and try the following. Each step is a live check on the production service.

What to do	What it proves
Open /api/eval	The gate re-runs: 55/55, grounding, abstain rate, no_auto_deny.
Open a case with a contradiction (CASE-2026-0003)	The verifier abstains; state is NEEDS_REVIEW, never DENIED.
Try to deny with an empty reason	The backend returns 400; a human act requires a reason.
Approve a case, then reopen its audit trail	The officer, reason, timestamp and hash are recorded.
Press 'Verify audit chain'	The chain recomputes to an intact head.
Press 'Tamper demo'	A forged copy breaks at the changed entry; the real ledger stays intact.
Open 'System & Adapters'	Requested vs active backends are reported honestly.

The design claim is deliberately modest and testable: the machine may propose, an independent check may refuse, and only a named person may decide. Everything else in Witness exists to make that sentence true in code, visible in the interface, and provable from the ledger.

Live service: https://alternate-echo-whale-5578-github-app.alternatefutures.ai

Repository: https://github.com/anithepani/GlueSticks-agentika